

SHIKSHAK
EDUCATION

EDUCATIONAL TECHNOLOGY

DIGITAL SHIKSHAK

SYSTEM ARCHITECTURE DOCUMENT

AUTHORS

Meet Bikhani

Lead Architect

Ayush Pandey

Educational Technology

Tanshik Sharma

AI Systems Engineer

Mausam Kar

Cloud Infrastructure

SHIKSHAK

System Architecture & Technical Specification

Document Information

Version: 1.0
Date: January 18, 2026
Status: Official Documentation
Classification: Technical Reference

About This Document

This comprehensive technical specification details the architecture, design philosophy, and implementation strategy of the SHIKSHAK Learning Management System. It serves as the authoritative reference for developers, architects, and stakeholders involved in the platform's development and deployment.

Target Audience

- Software Architects and System Designers
- Development Teams and Engineers
- DevOps and Infrastructure Teams
- Technical Product Managers
- Educational Technology Stakeholders

CHAPTER 0

Contents



1	Introduction to Shikshak	8
1.1	The Educational Revolution	8
1.2	The Problem We're Solving	9
1.2.1	Student Engagement Crisis	9
1.2.2	Academic Integrity	9
1.2.3	Teacher Burnout	9
1.2.4	Infrastructure Limitations	9
1.3	The Shikshak Solution	10
1.3.1	Intelligent Assistance	10
1.3.2	Verified Integrity	10

CONTENTS

1.3.3	Infinite Scalability	10
2	Target Audience & Use Cases	12
2.1	Higher Education Institutions	12
2.2	Corporate Training Departments	13
2.3	Coding Bootcamps & Professional Development	13
3	Architectural Philosophy	14
3.1	Microservices-First Design	14
3.1.1	Benefits of Microservices	14
3.2	Event-Driven Communication	15
3.3	Polyglot Persistence	16
4	Core Microservices	17
4.1	API Gateway Service	17
4.1.1	Purpose & Responsibilities	17
4.1.2	Request Flow	18
4.1.3	Technology Choices	18
4.2	Authentication Service	19
4.2.1	The Challenge of Identity	19
4.2.2	Multi-Strategy Authentication	19
4.2.3	Role-Based Access Control	19
4.2.4	Session Management	20
4.3	Course Management Service	20
4.3.1	Educational Content Hierarchy	20
4.3.2	Content Delivery	21
4.3.3	Progress Tracking	21
4.3.4	Dynamic Content Rendering	22
4.4	RAG (Retrieval-Augmented Generation) Service	22
4.4.1	The AI Brain of Shikshak	22
4.4.2	Understanding Vector Embeddings	22
4.4.3	The Indexing Process	23

4.4.4	The Query Process	23
4.4.5	Preventing Hallucinations	24
4.5	Payment Service	24
4.5.1	Transactional Integrity	24
4.5.2	The Purchase Flow	25
4.5.3	Revenue Distribution	25
4.6	Notification Service	26
4.6.1	Keeping Users Informed	26
4.6.2	Notification Types	26
4.6.3	Smart Batching	26
4.6.4	Preference Management	27
5	Technology Stack Deep Dive	28
5.1	Frontend: Next.js Framework	28
5.1.1	Why Next.js	28
5.1.2	State Management with Zustand	29
5.1.3	Styling with Tailwind CSS	29
5.2	Backend: Node.js Ecosystem	29
5.2.1	Non-Blocking I/O	29
5.2.2	npm Package Ecosystem	30
5.3	Database: MongoDB	30
5.3.1	Document-Oriented Schema	30
5.3.2	Horizontal Scaling with Sharding	30
5.3.3	Replica Sets for High Availability	30
5.4	Message Queue: Apache Kafka	30
5.4.1	What Makes Kafka Special	31
5.4.2	Topics and Partitions	31
5.4.3	Event Replay	31
5.5	Caching & Vector Storage: Redis	31
5.5.1	Microsecond Latency	32
5.5.2	Vector Search Capabilities	32

CONTENTS

5.6	Cloud Infrastructure: Azure	32
5.6.1	Blob Storage	32
5.6.2	OpenAI Service	32
6	Advanced Features	34
6.1	AI Proctoring System	34
6.1.1	The Challenge of Remote Assessment Integrity	34
6.1.2	Face Verification	34
6.1.3	Gaze Tracking	35
6.1.4	Environment Monitoring	35
6.1.5	Privacy-First Design	36
6.2	Adaptive Learning Paths	36
6.2.1	Personalized Progression	36
6.2.2	Learning Analytics	37
6.3	Collaborative Features	37
6.3.1	Discussion Forums	37
6.3.2	Study Groups	37
6.3.3	Peer Review	38
7	Short-Term Enhancements (6-12 Months)	39
7.1	Mobile Application	39
7.2	Enhanced Analytics Dashboard	40
7.3	Automated Content Generation	40
7.4	Accessibility Improvements	40
8	Medium-Term Innovations (1-2 Years)	41
8.1	Virtual Labs & Simulations	41
8.2	Live Virtual Classrooms	42
8.3	Blockchain Certificates	42
8.4	Advanced AI Tutoring	42
8.5	Institutional Integration	43

9 Long-Term Vision (3-5 Years)	44
9.1 Personalized AI Agents	44
9.2 VR/AR Immersive Learning	44
9.3 Federated Learning for Privacy	45
9.4 Global Education Network	45
9.5 Neurofeedback Integration	45
10 Conclusion	47
10.1 The Path Forward	47
10.2 Open Challenges	48
10.3 Impact Potential	48

CHAPTER 1

Introduction to Shikshak

1.1 The Educational Revolution

The world of education stands at a crossroads. Traditional Learning Management Systems have served us well for decades, but they were designed for an era of passive content consumption. Students watch videos, read PDFs, and take quizzes—but the system doesn't truly understand them. It doesn't adapt to their learning pace, doesn't answer their questions in real-time, and certainly doesn't ensure academic integrity in remote environments.

Shikshak represents a fundamental reimagining of what an LMS can be. We're not just adding AI features to an existing platform; we're building an **Intelligent Educational Ecosystem** from the ground up, where every architectural decision is made with artificial intelligence, scalability, and academic integrity as core principles.



Figure 1.1: Core Pillars of the SHIKSHAK Platform

1.2 The Problem We're Solving

Modern education faces several critical challenges that demand innovative solutions:

1.2.1 Student Engagement Crisis

Studies show that online course completion rates hover around 5-15%. Students feel isolated, confused, and unsupported. When they have questions at midnight while studying, they have nowhere to turn except generic search engines that don't understand their specific course context.

1.2.2 Academic Integrity

The shift to remote learning exposed a fundamental weakness: we have no reliable way to verify that students are who they claim to be during assessments, and that they're completing work independently.

1.2.3 Teacher Burnout

Instructors are overwhelmed answering the same questions repeatedly, manually grading assessments, and trying to identify struggling students before it's too late.

1.2.4 Infrastructure Limitations

Traditional LMS platforms struggle when thousands of students attempt to access content simultaneously—especially during exam periods or assignment deadlines.

1.3 The Shikshak Solution

Our platform addresses these challenges through three revolutionary pillars:

1.3.1 Intelligent Assistance

Every student has access to an **AI Study Companion** that has ingested and understood their entire course material. This isn't a generic chatbot—it's a specialized tutor that can explain complex concepts, provide worked examples, and guide students through problem-solving processes, all while being strictly grounded in the course content to prevent misinformation.

AI Study Companion Features

- 24/7 availability for instant question answering
- Context-aware responses based on course material
- Step-by-step problem-solving guidance
- Multimodal understanding (text, images, code)
- Personalized learning recommendations

1.3.2 Verified Integrity

Our browser-based **AI Proctoring Suite** uses computer vision to maintain assessment integrity without invasive software installation. It monitors student identity, environment, and behavior patterns in real-time, flagging anomalies for instructor review while respecting privacy by processing most data client-side.

1.3.3 Infinite Scalability

Built on Event-Driven Microservices Architecture, Shikshak can handle 10 students or 100,000 with equal ease. The system automatically scales individual services based on demand, ensuring teachers never have to worry about whether the infrastructure can handle their course enrollment.

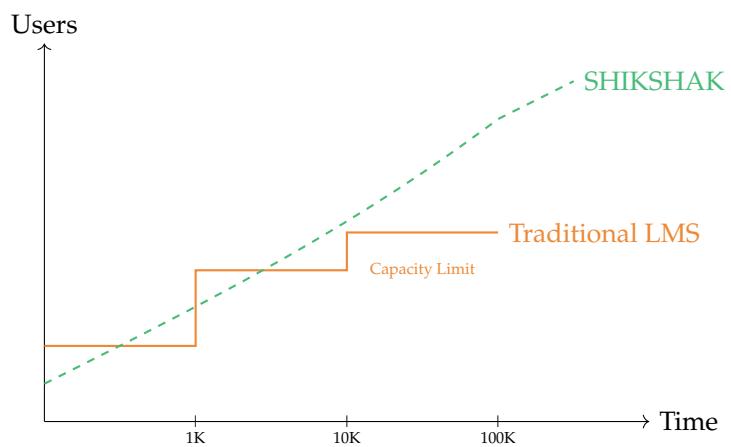


Figure 1.2: Scalability Comparison: Traditional LMS vs. SHIKSHAK

CHAPTER 2

Target Audience & Use Cases



2.1 Higher Education Institutions

Universities and colleges face increasing pressure to offer robust online and hybrid learning options. Shikshak provides:

- **MOOC-Scale Support:** Handle massive open online courses with thousands of concurrent students, each receiving personalized AI assistance
- **Credible Certifications:** Issue certificates and degrees with confidence, knowing that assessments were completed with verified integrity
- **Analytics for Accreditation:** Comprehensive learning analytics help demonstrate educational outcomes for accreditation bodies

2.2 Corporate Training Departments

Enterprises investing in employee development need systems that deliver measurable results:

- **Completion Verification:** Prove that employees actually completed required training modules with proctored assessments
- **Just-In-Time Learning:** AI assistants help employees find specific information within training materials instantly
- **Cost Efficiency:** Reduce dependency on live instructors by automating routine questions and grading

2.3 Coding Bootcamps & Professional Development

Skills-based education providers benefit from:

- **Project-Based Assessment:** Monitor students as they complete real-world projects, ensuring independent work
- **Technical Deep-Dives:** AI can explain complex programming concepts, debug logic errors conceptually, and guide problem-solving
- **Career Credibility:** Employers trust certifications backed by verified assessment integrity

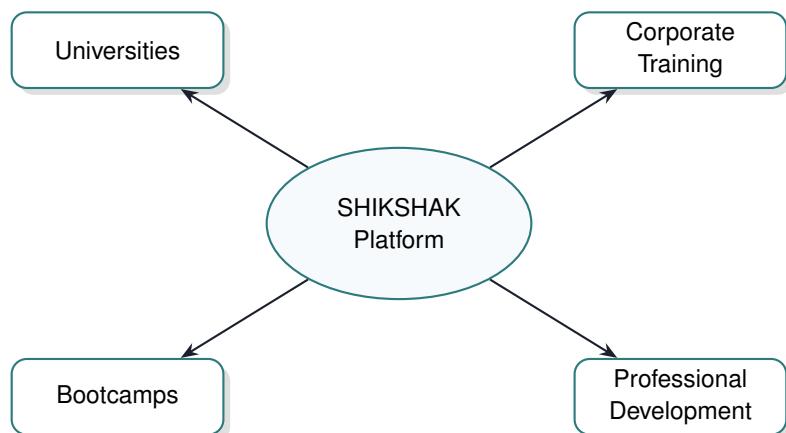


Figure 2.1: SHIKSHAK's Diverse Application Domains

CHAPTER 3

Architectural Philosophy



3.1 Microservices-First Design

Shikshak is built entirely on microservices architecture. This isn't just a trendy buzzword—it's a fundamental architectural choice that provides concrete benefits:

Key Architectural Principle

Each service is independently deployable, scalable, and maintainable. This separation of concerns allows teams to work in parallel and scale resources precisely where needed.

3.1.1 Benefits of Microservices

Independent Scaling: The AI RAG service requires significant computational resources and may need 10 server instances during peak hours. Meanwhile, the Authentication

service is lightweight and might only need 2 instances. Microservices allow us to scale each component independently based on its specific demands.

Technology Flexibility: Each service can be built with the technology best suited to its purpose. The RAG service might use Python for its rich AI/ML ecosystem, while the Course Service uses Node.js for fast I/O operations. They communicate via standard HTTP and Kafka protocols.

Fault Isolation: If the Payment Service experiences a bug and crashes, students can still watch videos, ask AI questions, and submit assignments. The system degrades gracefully rather than failing catastrophically.

Team Autonomy: Different development teams can own different services, working in parallel without stepping on each other's toes. The Course Service team can deploy updates without coordinating with the RAG Service team.



Figure 3.1: Evolution from Monolithic to Microservices Architecture

3.2 Event-Driven Communication

Not all operations need immediate responses. When a teacher uploads a 2GB video file, we don't want them waiting 15 minutes while our servers transcode it into streaming formats. Instead, we use an asynchronous, event-driven approach:

1. The upload completes in seconds, and the teacher receives confirmation
2. An event is published to Apache Kafka: "New video uploaded for Course XYZ"
3. A background worker service picks up this event and begins processing
4. When complete, another event triggers: "Video processing complete for Course XYZ"
5. The Course Service updates the database, and a notification is sent to the teacher

This asynchronous pattern keeps the user interface responsive and allows heavy operations to happen in the background without blocking user interactions.

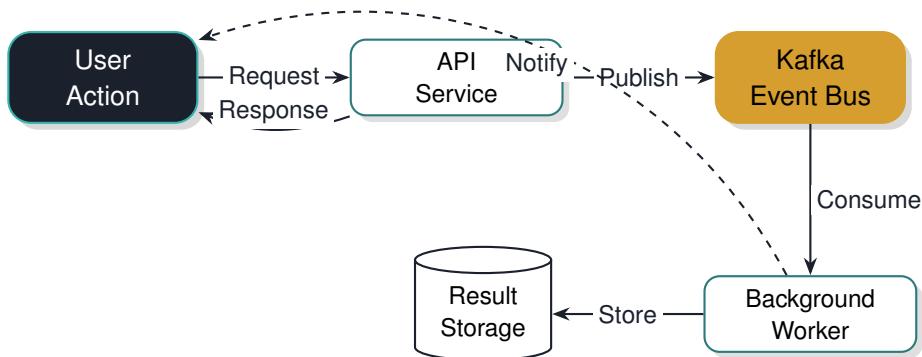


Figure 3.2: Event-Driven Asynchronous Processing Pattern

3.3 Polyglot Persistence

Different types of data have different storage requirements. Shikshak uses the right database for each job:

- **MongoDB:** Stores hierarchical course structures (Courses → Modules → Lessons) with flexible schemas that can evolve as educational requirements change
- **Redis:** Serves dual purposes—ultra-fast caching for frequently accessed data, and vector storage for AI embeddings that power semantic search
- **Azure Blob Storage:** Handles large binary files like videos and documents efficiently with CDN integration for global content delivery

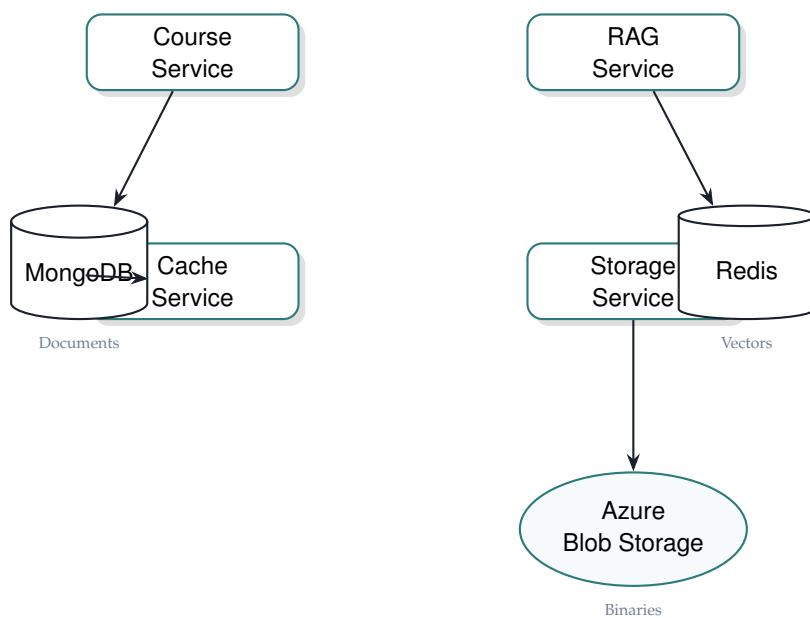


Figure 3.3: Polyglot Persistence Strategy

CHAPTER 4

Core Microservices



4.1 API Gateway Service

4.1.1 Purpose & Responsibilities

The Gateway is the front door to our entire backend infrastructure. Every HTTP request from the client must pass through this service. This centralization provides several critical advantages:

- **Security Checkpoint:** SSL certificates are managed here. All HTTPS traffic is decrypted at the gateway, inspected for threats, and then forwarded to internal services over a trusted private network
- **Unified Entry Point:** Clients only need to know one URL. The gateway handles routing requests to the appropriate microservice based on the path
- **Rate Limiting:** Prevents abuse by limiting how many requests a single user can make per

minute, protecting backend services from being overwhelmed

- **Request/Response Transformation:** Can modify headers, add authentication tokens, or transform data formats as needed

4.1.2 Request Flow

When a student's browser sends a request to view "Course 123, Module 5, Lesson 2", the Gateway:

1. Receives the request at: /api/v1/courses/123/modules/5/lessons/2
2. Validates the JWT token in the request header to confirm the user is logged in
3. Checks rate limits: Has this user made too many requests recently?
4. Routes the request to the Course Service's internal address
5. Receives the response from Course Service
6. Adds CORS headers if needed for browser compatibility
7. Returns the response to the client's browser

All of this happens in milliseconds, providing a seamless experience while maintaining security and organization.

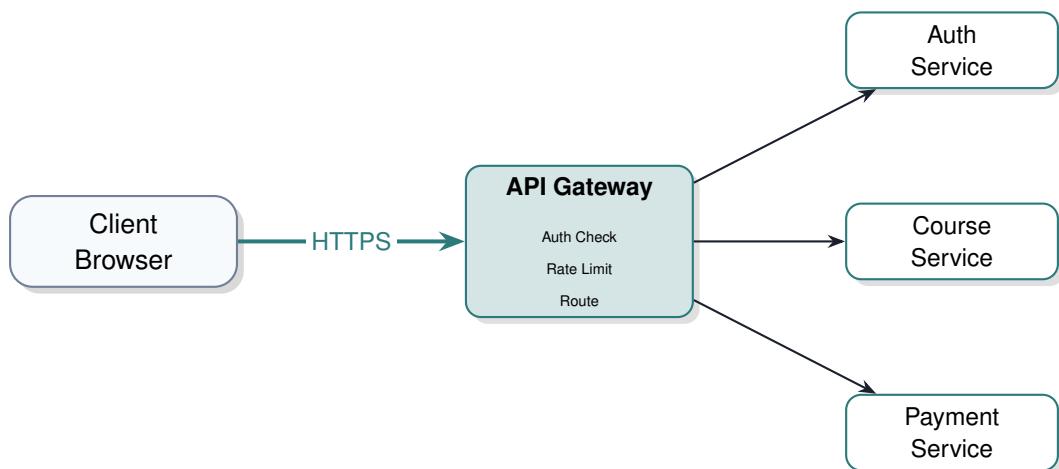


Figure 4.1: API Gateway Request Routing

4.1.3 Technology Choices

The Gateway is built with **Node.js** using the **Express** framework, chosen for its non-blocking I/O model that allows it to handle thousands of concurrent connections efficiently. It uses minimal CPU because it's primarily forwarding requests rather than processing them.

4.2 Authentication Service

4.2.1 The Challenge of Identity

In educational systems, identity is paramount. We need to answer three questions with certainty:

- **Authentication:** Is this user who they claim to be?
- **Authorization:** What is this user allowed to access?
- **Accountability:** Can we maintain an audit trail of user actions?

4.2.2 Multi-Strategy Authentication

Shikshak supports multiple ways for users to prove their identity:

OAuth 2.0 Integration: Users can sign in with Google or GitHub. When they click "Sign in with Google," we redirect them to Google's servers. Google verifies their identity (username, password, two-factor authentication), then sends us back an authorization code confirming their identity. We never see or store their Google password.

Email & Password: Traditional authentication for users who prefer it. Passwords are hashed using bcrypt with a computational cost factor of 12, meaning even if our database is compromised, the passwords remain computationally infeasible to crack.

Magic Links: For quick access, users can request a one-time login link sent to their verified email address. The link contains a cryptographically secure token valid for 15 minutes.

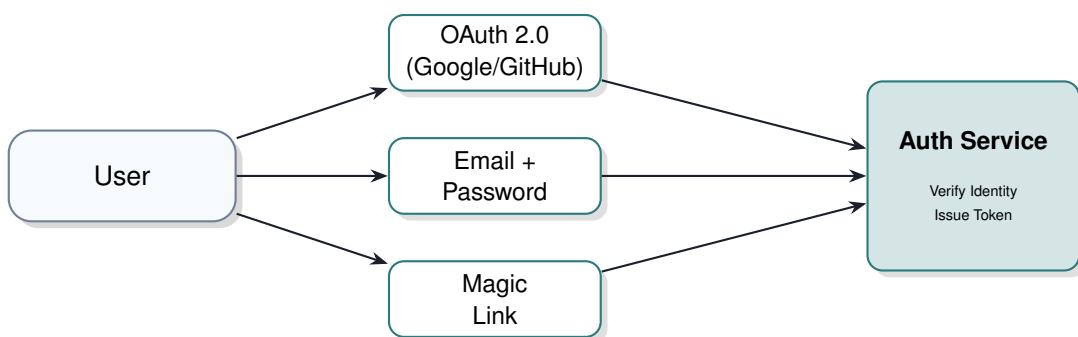


Figure 4.2: Multi-Strategy Authentication Flow

4.2.3 Role-Based Access Control

Not all users are equal. Shikshak defines several roles:

- **Students:** Can view courses they're enrolled in, submit assignments, and take quizzes
- **Teachers:** Can create and modify their own courses, view student progress, and grade assessments

- **Administrators:** Can manage users, configure system settings, and access platform analytics
- **Institution Admins:** Have elevated privileges for their specific organization, managing teachers and students within their institution

The Auth Service maintains a mapping of users to roles and enforces these permissions through middleware that intercepts every request.

4.2.4 Session Management

When you log in, the Auth Service creates a session token—a random, cryptographically secure string. This token is stored:

1. In the database, associated with your user ID and an expiration timestamp
2. In your browser as an HTTP-only cookie (JavaScript cannot access it, preventing XSS attacks)

Every subsequent request includes this cookie. The Auth Service validates it by checking if the token exists in the database and hasn't expired. Sessions expire after 7 days of inactivity for security.

4.3 Course Management Service

4.3.1 Educational Content Hierarchy

Online education has a natural structure that Shikshak mirrors in its data model:

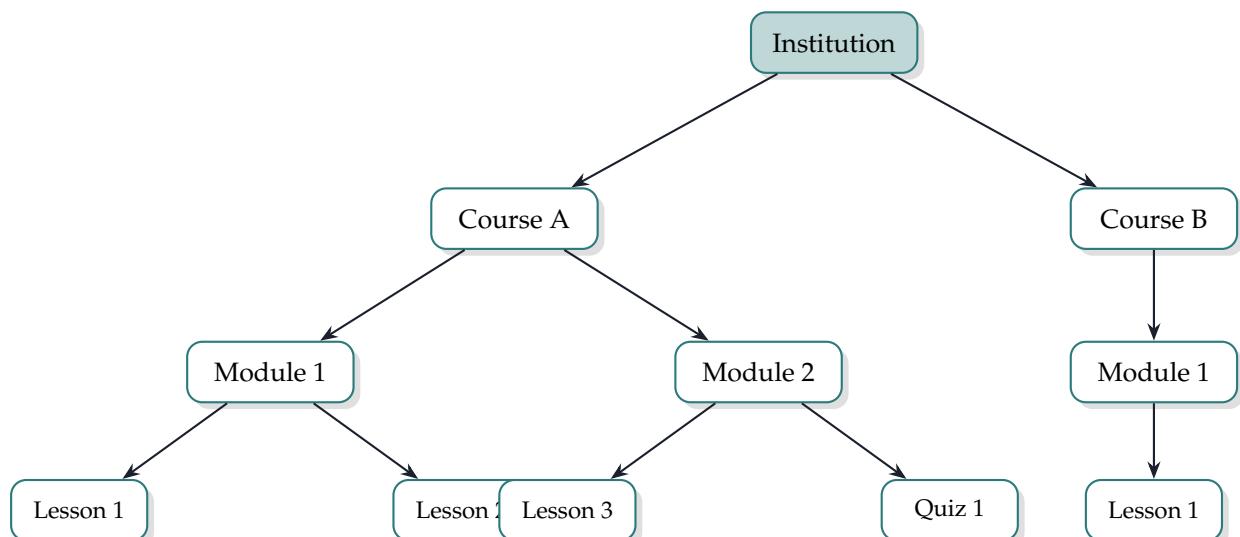


Figure 4.3: Educational Content Hierarchy

Institution Level: A university or company running the platform.

Course Level: An individual class, like "Introduction to Machine Learning" or "Corporate Ethics Training."

Module Level: Major topics within a course, such as "Neural Networks" or "Decision Trees."

Lesson Level: Individual learning units, typically a video lecture, reading, or interactive activity.

Assessment Level: Quizzes, exams, and assignments that test understanding.

The Course Service manages this entire hierarchy, tracking relationships and permissions.

4.3.2 Content Delivery

When a student clicks to watch a video lesson, the Course Service:

1. Verifies the student is enrolled in the course (authorization check)
2. Checks if any prerequisite lessons must be completed first
3. Retrieves the video URL from cloud storage
4. Generates a time-limited, signed URL valid for 2 hours (prevents sharing links)
5. Logs that the student accessed this lesson (for progress tracking)
6. Returns the URL along with metadata (title, description, duration, transcript)

4.3.3 Progress Tracking

The service maintains detailed records of student progress:

- Which lessons have been viewed and what percentage was watched
- Quiz scores and attempt timestamps
- Assignment submission status
- Time spent on each lesson (captured via heartbeat pings from the video player)
- Completion status for modules and overall course

Teachers can access these analytics to identify struggling students and adapt their teaching strategy.

Progress Metrics

The Course Service tracks over 20 different engagement metrics per student, including:

- Video completion rates

- Quiz attempt patterns
- Time-on-task measurements
- Discussion forum participation
- Assignment submission timeliness

4.3.4 Dynamic Content Rendering

Course content isn't static. Teachers can:

- Reorder modules and lessons by dragging them in the UI
- Hide lessons until a specific date (for phased content release)
- Update lesson content and have changes reflect immediately for all students
- Attach supplementary materials (PDFs, slides, external links)

The Course Service handles all this complexity while maintaining performance even with thousands of students accessing content simultaneously.

4.4 RAG (Retrieval-Augmented Generation) Service

4.4.1 The AI Brain of Shikshak

This is perhaps the most technically sophisticated service in the entire platform. It's what enables students to ask natural language questions and receive accurate, context-aware answers grounded in their course material.

4.4.2 Understanding Vector Embeddings

Computers don't understand language the way humans do. They work with numbers. Vector embeddings are a way to convert text into high-dimensional numerical coordinates that capture semantic meaning.

For example, the sentence "The cat sat on the mat" might become a vector like:

[0.23, -0.45, 0.89, ..., 0.12]

(with typically 768 or 1536 dimensions).

Similar meanings produce similar vectors. "The feline rested on the rug" would have a vector mathematically close to the first one, even though the exact words are different. This allows us to search by meaning rather than just keyword matching.

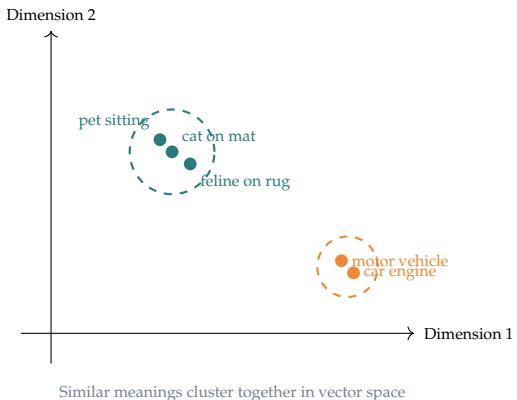


Figure 4.4: Conceptual Visualization of Vector Embeddings

4.4.3 The Indexing Process

When a teacher uploads course material (video transcripts, lecture notes, textbook chapters), the RAG Service:

1. **Chunks the Content:** Breaks large documents into manageable pieces (typically 500-1000 words each with 100-word overlap to maintain context across boundaries)
2. **Generates Embeddings:** Each chunk is sent to an embedding model that converts the text into a vector. Shikshak uses OpenAI's text-embedding-ada-002 model or Azure's equivalent
3. **Stores in Vector Database:** The vectors are stored in Redis with its vector search capabilities, indexed for fast similarity lookups
4. **Maintains Metadata:** Each vector is tagged with metadata: which course it belongs to, what module, lesson, page number, etc.

4.4.4 The Query Process

When a student asks "*How does gradient descent work in backpropagation?*", the RAG Service:

1. **Vectorizes the Question:** The question is converted into a vector using the same embedding model
2. **Searches for Similar Content:** Redis performs a cosine similarity search, finding the top 5-10 chunks from the course material that are most semantically similar to the question vector
3. **Constructs a Prompt:** The retrieved chunks are assembled into a prompt for the LLM
4. **Generates Response:** The prompt is sent to Azure OpenAI (GPT-4) which generates a comprehensive answer grounded in the provided material
5. **Returns with Citations:** The answer is sent back to the student, often with references to specific lessons where they can learn more

This entire process completes in 2-4 seconds, providing students with instant, accurate assistance 24/7.

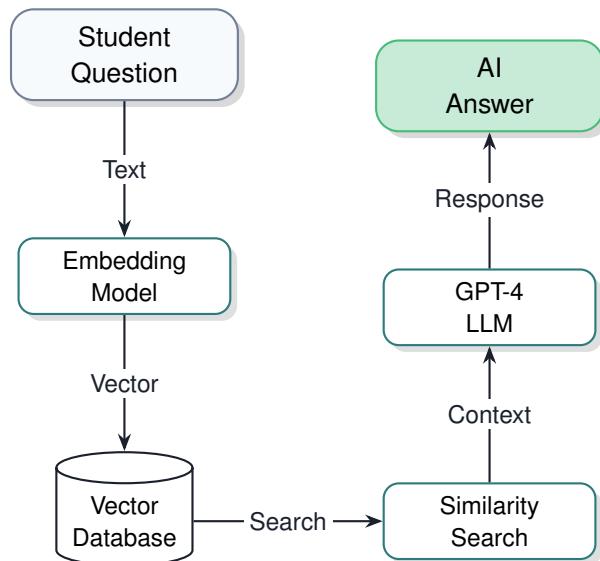


Figure 4.5: RAG Service Query Flow

4.4.5 Preventing Hallucinations

One of the biggest challenges with AI is "hallucination"—when models generate false information confidently. The RAG approach mitigates this by:

- Constraining the AI to only answer based on provided course material
- Including explicit instructions to admit when it doesn't know something
- Showing students which course sections the answer was derived from
- Allowing teachers to review and override AI responses

Quality Control

All AI-generated responses include confidence scores and source citations. Teachers can flag incorrect responses, which are used to continuously improve the system through fine-tuning and prompt engineering.

4.5 Payment Service

4.5.1 Transactional Integrity

Handling payments in an educational platform requires absolute reliability. Students must never be charged twice, and teachers must always receive credit for their course sales. The

Payment Service ensures this through careful transaction management.

4.5.2 The Purchase Flow

When a student clicks "Enroll" on a paid course:

1. **Price Verification:** The service confirms the current price (in case the teacher changed it since the student opened the page)
2. **Duplicate Check:** Verifies the student hasn't already purchased this course
3. **Payment Gateway:** Creates a checkout session with Stripe or equivalent, generating a secure payment page
4. **Webhook Handling:** When Stripe confirms payment, a webhook notifies our Payment Service immediately
5. **Database Update:** The student is marked as enrolled in the course atomically (either the entire operation succeeds or none of it does)
6. **Event Publication:** A Kafka event announces "Student X enrolled in Course Y" so other services can react (send welcome email, update analytics, etc.)

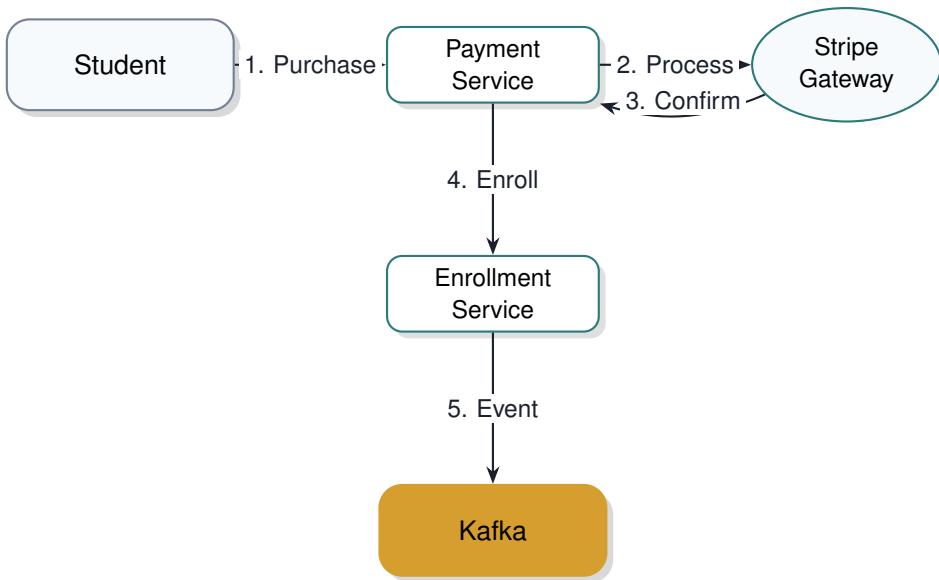


Figure 4.6: Secure Payment Processing Flow

4.5.3 Revenue Distribution

For marketplace-style deployments where multiple teachers sell courses on a shared platform, the Payment Service handles complex revenue splitting:

- Platform fee (e.g., 20% to Shikshak operators)

- Teacher commission (e.g., 75% to course creator)
- Affiliate commission (e.g., 5% to the blog that referred the student)

This is all configured per course and calculated automatically at purchase time.

4.6 Notification Service

4.6.1 Keeping Users Informed

An LMS is ineffective if students miss important updates. The Notification Service ensures timely, relevant communication across multiple channels.

4.6.2 Notification Types

Email Notifications: Sent for major events like course enrollment confirmation, assignment grading, or deadline reminders. Uses SendGrid or AWS SES for reliable delivery with bounce handling.

In-App Notifications: Displayed in the web interface with a bell icon counter. Students see these immediately when logged in.

Push Notifications: For the mobile app (future feature), delivered via Firebase Cloud Messaging.

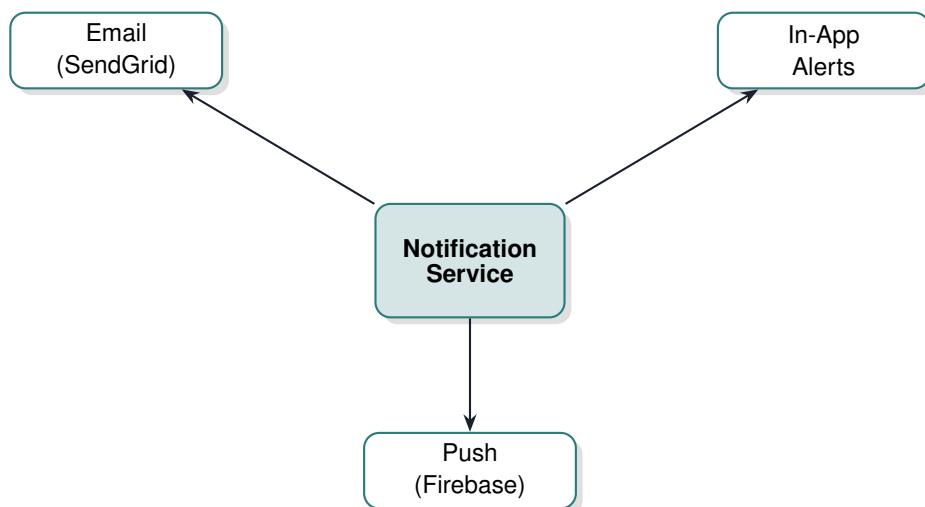


Figure 4.7: Multi-Channel Notification System

4.6.3 Smart Batching

Instead of sending 50 separate emails if a teacher posts 50 new announcements, the Notification Service batches them into a single daily digest email, reducing inbox overload while ensuring students stay informed.

4.6.4 Preference Management

Students can customize which notifications they receive, opting out of non-essential updates while still receiving critical information like grade releases and deadline reminders.

CHAPTER 5

Technology Stack Deep Dive



5.1 Frontend: Next.js Framework

5.1.1 Why Next.js

Shikshak's frontend is built with **Next.js 15**, a React-based framework that provides several critical advantages for an LMS:

- **Server-Side Rendering (SSR):** When a student clicks a link to a course, the server pre-renders the HTML with all the course information before sending it to the browser. This provides instant content visibility and excellent SEO
- **Static Site Generation (SSG):** For content that doesn't change often (like the marketing homepage), Next.js can generate HTML at build time, serving it with blazing speed from a CDN
- **App Router:** The latest Next.js architecture organizes code by feature, making large

applications more maintainable

- **Image Optimization:** Automatically resizes and optimizes course thumbnails and profile pictures

5.1.2 State Management with Zustand

While Next.js handles server state beautifully, client-side state (like "is the video player currently playing" or "what's the user's selected theme") requires dedicated management. Shikshak uses **Zustand**, a lightweight state library that's simpler than Redux but more powerful than React's built-in Context API.

5.1.3 Styling with Tailwind CSS

Every UI component is styled using **Tailwind CSS 4**'s utility-first approach. Instead of writing custom CSS classes, developers apply pre-built classes directly in the HTML. This ensures design consistency and speeds up development significantly.

5.2 Backend: Node.js Ecosystem

5.2.1 Non-Blocking I/O

Most LMS operations are I/O-bound rather than CPU-bound. When fetching data from a database or calling an external API, the server spends most of its time waiting for responses. Node.js handles this brilliantly with its event loop architecture.

Traditional threaded servers create a new thread for each request, which consumes memory. Node.js uses a single thread with asynchronous callbacks, allowing thousands of concurrent connections with minimal resource usage.

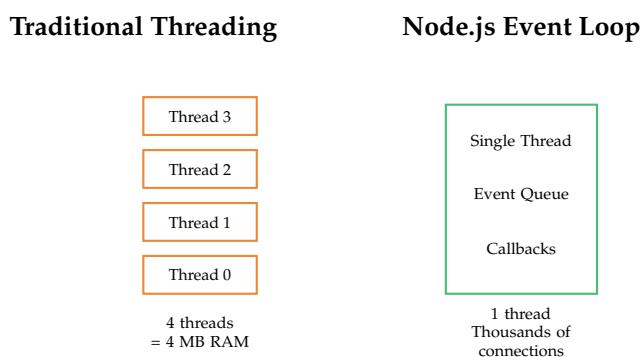


Figure 5.1: Threading Model Comparison

5.2.2 npm Package Ecosystem

Node.js has the world's largest package registry. Need JWT authentication? Install `jsonwebtoken`. Need to resize images? Install `sharp`. This rich ecosystem allows rapid development without reinventing wheels.

5.3 Database: MongoDB

5.3.1 Document-Oriented Schema

Educational content is hierarchical and variable in structure. A Biology course might have lab videos, while a Mathematics course has interactive problem sets. MongoDB's flexible document schema adapts perfectly.

5.3.2 Horizontal Scaling with Sharding

When Shikshak grows to millions of users, a single database server won't suffice. MongoDB supports sharding—splitting data across multiple servers based on a shard key (like user ID or course ID). Queries automatically route to the correct shard, distributing load.

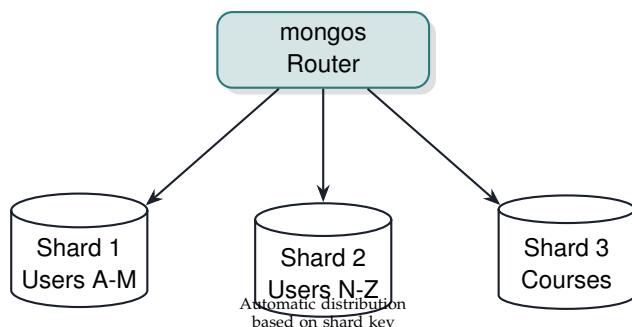


Figure 5.2: MongoDB Sharding Architecture

5.3.3 Replica Sets for High Availability

Every production MongoDB deployment runs with replica sets: one primary server handles writes, while multiple secondary servers replicate data in real-time. If the primary fails, a secondary automatically promotes itself to primary within seconds, ensuring near-zero downtime.

5.4 Message Queue: Apache Kafka

5.4.1 What Makes Kafka Special

Kafka isn't just a message queue—it's a distributed streaming platform built for massive scale. While traditional queues like RabbitMQ work well for moderate loads, Kafka can handle millions of events per second.

5.4.2 Topics and Partitions

Kafka organizes events into topics. Shikshak might have topics like:

- `user.registered` - New user signups
- `video.uploaded` - New course videos
- `assessment.submitted` - Student quiz/exam submissions
- `payment.completed` - Successful course purchases

Each topic is further divided into partitions for parallel processing. If video processing workers need to handle high load, they can consume from multiple partitions simultaneously.



Figure 5.3: Kafka Event Streaming Architecture

5.4.3 Event Replay

Unlike traditional queues that delete messages after consumption, Kafka retains events for a configured period (e.g., 7 days). This allows:

- **Debugging:** If a bug caused incorrect processing last week, we can replay those events through a fixed version of the code
- **Analytics:** A new analytics service can read historical events to build dashboards retroactively
- **Disaster Recovery:** If a database is corrupted, we can rebuild it by replaying all events in order

5.5 Caching & Vector Storage: Redis

5.5.1 Microsecond Latency

Redis stores data entirely in RAM, providing sub-millisecond response times. This makes it perfect for:

- **Session Storage:** User session tokens are stored in Redis for instant authentication checks
- **Leaderboard Caching:** Course leaderboards showing top-performing students are cached for 5 minutes rather than recalculating from the database on every page load
- **Rate Limiting:** Track how many API requests each user has made in the last minute, enforcing limits without database overhead

5.5.2 Vector Search Capabilities

Recent Redis versions include RedisSearch with vector similarity searching. This allows the RAG service to store and query millions of embedding vectors efficiently without requiring a specialized vector database like Pinecone or Weaviate.

5.6 Cloud Infrastructure: Azure

5.6.1 Blob Storage

All video files, PDFs, and user-uploaded content live in **Azure Blob Storage**. Benefits include:

- **Scalability:** Store petabytes of data without managing disk arrays
- **CDN Integration:** Integrate with Azure CDN for global content delivery with servers near students worldwide
- **Lifecycle Policies:** Automatically move old content to cheaper "cool" or "archive" tiers

5.6.2 OpenAI Service

Azure's OpenAI deployment provides enterprise-grade access to GPT-4 with guaranteed uptime SLAs, data residency compliance, and usage quotas that ensure consistent RAG service performance.

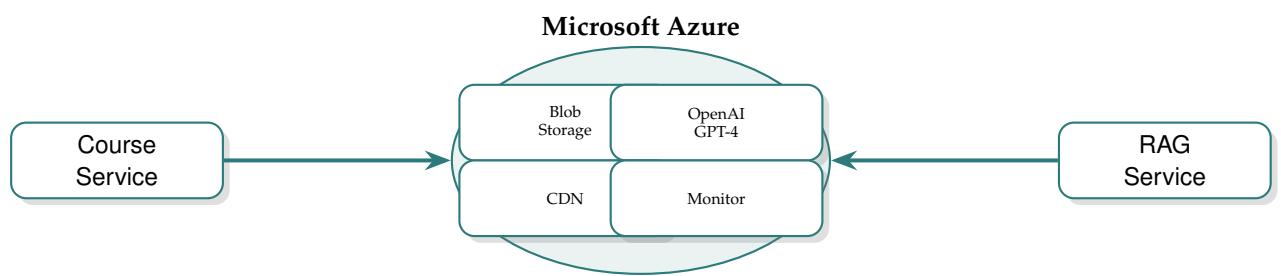


Figure 5.4: Azure Cloud Services Integration

CHAPTER 6

Advanced Features



6.1 AI Proctoring System

6.1.1 The Challenge of Remote Assessment Integrity

Traditional in-person exams provide natural safeguards: an instructor watches students, controls the environment, and verifies identities. Remote exams lack these controls, creating opportunities for cheating. Shikshak's AI Proctoring System restores integrity while respecting student privacy.

6.1.2 Face Verification

Before an exam begins, students complete an enrollment process:

1. **ID Capture:** Student uploads a photo of their government ID

2. **Live Photo:** Student takes a selfie using their webcam
3. **Face Encoding:** Our system (using MediaPipe's Face Detection) creates a mathematical representation of facial features—not storing the actual photo to protect privacy

During the exam:

1. Webcam captures periodic snapshots (every 60 seconds)
2. Faces are detected and encoded client-side (in the browser)
3. Encodings are compared to the enrolled face
4. If no face is detected or a different person appears, an alert is flagged for instructor review

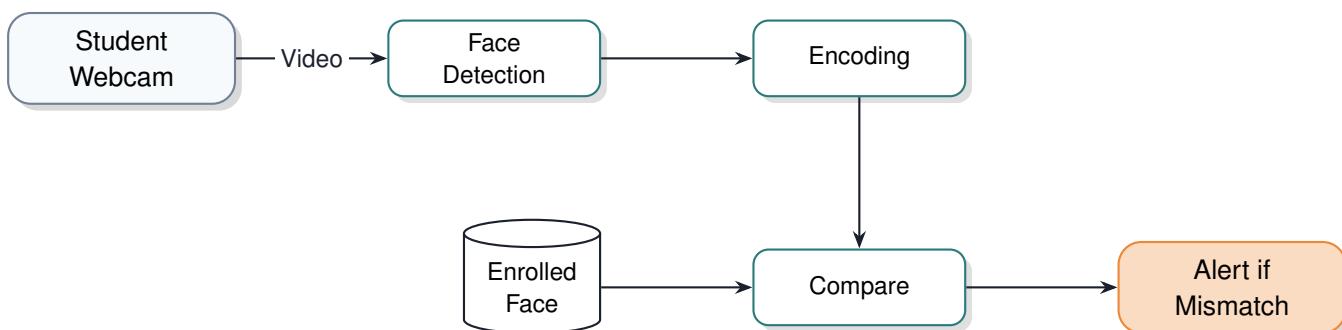


Figure 6.1: Face Verification Pipeline

6.1.3 Gaze Tracking

Students looking away from the screen for extended periods may be reading notes or using a second device. The system tracks:

- Eye position relative to screen (using pupil detection)
- Head pose (looking left, right, up, down)
- Duration of off-screen attention

Flags are raised if students look away for more than 10 seconds cumulatively per question, but the system accounts for normal glances around to think.

6.1.4 Environment Monitoring

The proctoring system analyzes the exam environment:

- **Person Detection:** If multiple faces appear in the webcam feed, it suggests someone is helping the student

- **Audio Monitoring:** Detects if the student is speaking (possibly communicating with someone off-screen), though actual audio content is not recorded for privacy
- **Tab Switching:** Browser JavaScript detects when students switch to other tabs or applications, which might indicate searching for answers online

6.1.5 Privacy-First Design

Unlike traditional proctoring software that records entire exam sessions and uploads gigabytes of video, Shikshak:

- Processes most data client-side (in the browser) using TensorFlow.js
- Only sends violation events (specific frames where issues were detected)
- Deletes all proctoring data 30 days after exam completion
- Gives students transparency into exactly what is monitored
- Allows students to request deletion of their biometric data

Privacy Commitment

SHIKSHAK is committed to student privacy. All biometric data is processed locally when possible, and stored data is encrypted and automatically purged after 30 days. Students have full transparency and control over their data.

6.2 Adaptive Learning Paths

6.2.1 Personalized Progression

Not all students learn at the same pace. Shikshak's adaptive system adjusts content delivery based on performance:

- **Prerequisite Enforcement:** If a student struggles with "Module 3: Derivatives," the system might suggest reviewing "Module 2: Limits" before progressing
- **Difficulty Adjustment:** Quiz questions are selected from a bank, with difficulty adapting to the student's performance
- **Content Recommendations:** The system suggests supplementary materials based on areas where the student shows weakness

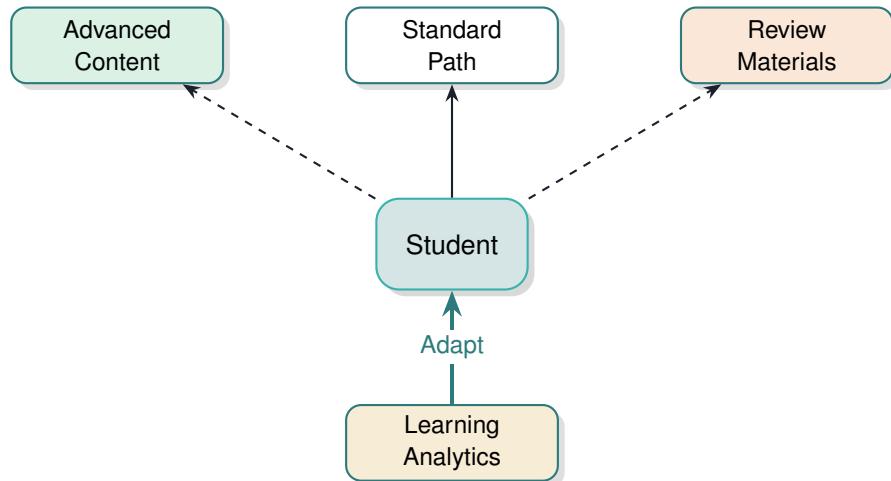


Figure 6.2: Adaptive Learning Path Selection

6.2.2 Learning Analytics

Teachers receive insights powered by machine learning:

- Which lessons cause the most students to rewatch content multiple times?
- What quiz questions have the lowest success rates (indicating unclear teaching)?
- Which students are at risk of falling behind based on engagement patterns?
- How does student performance correlate with time-of-day or day-of-week?

These insights drive continuous course improvement.

6.3 Collaborative Features

6.3.1 Discussion Forums

Each course includes threaded discussion boards where:

- Students ask questions and help each other
- Teachers can mark official answers
- The AI assistant can suggest relevant past discussions when students ask questions
- Participation contributes to student engagement scores

6.3.2 Study Groups

Students can form virtual study groups within a course:

- Shared notes and flashcard decks
- Group video calls (integrated with WebRTC)
- Collaborative whiteboards for working through problems together
- Scheduled study sessions with calendar integration

6.3.3 Peer Review

For essay assignments and projects, teachers can enable peer review:

1. Students submit their work by a deadline
2. The system randomly assigns each student 3 peers' submissions to review
3. Students provide feedback using a rubric
4. Teachers review the peer feedback and assign final grades

This teaches critical thinking and evaluation skills while lightening the teacher's grading burden.

CHAPTER 7

Short-Term Enhancements (6-12 Months)

7.1 Mobile Application

Building native iOS and Android apps using React Native to provide:

- Offline video playback (download lectures for viewing without internet)
- Push notifications for assignment reminders and grade releases
- Biometric authentication (Face ID, fingerprint)
- Optimized UI for smaller screens

7.2 Enhanced Analytics Dashboard

Expanding the instructor analytics with:

- **Predictive Analytics:** Machine learning models that predict which students are likely to drop out based on engagement patterns, allowing early intervention
- **Comparative Analysis:** Show how current course performance compares to previous semesters
- **Content Effectiveness:** Heat maps showing which video segments students rewatch most, indicating areas of confusion

7.3 Automated Content Generation

Using AI to assist teachers with course creation:

- Generate quiz questions automatically from lecture transcripts
- Create summary flashcards for each lesson
- Suggest structure and pacing for new courses based on similar successful courses

7.4 Accessibility Improvements

Making education truly inclusive:

- Automatic caption generation in multiple languages
- Audio descriptions for visual content (for visually impaired students)
- Screen reader optimization throughout the platform
- High-contrast themes for users with visual impairments
- Keyboard navigation for all features

CHAPTER 8

Medium-Term Innovations (1-2 Years)

8.1 Virtual Labs & Simulations

For STEM courses, integrate interactive simulations:

- **Physics Labs:** Simulate projectile motion, electricity circuits, or optics experiments in a virtual environment
- **Chemistry Labs:** Mix virtual chemicals and observe reactions without physical lab access
- **Programming Sandboxes:** Embedded code editors where students write and execute code directly in lessons, with automated test case verification

8.2 Live Virtual Classrooms

Real-time video conferencing optimized for education:

- Integrated whiteboard with math equation support
- Screen sharing with annotation tools
- Breakout rooms for small group activities
- Recording and automatic transcription for students who miss live sessions
- Polls and quizzes during live sessions to check understanding

8.3 Blockchain Certificates

Issue tamper-proof credentials:

- Course completion certificates stored on a blockchain
- Students own their credentials, which can be verified by employers instantly
- Micro-credentials for completing individual modules or skill pathways
- Stackable credentials that combine into larger certifications or degrees

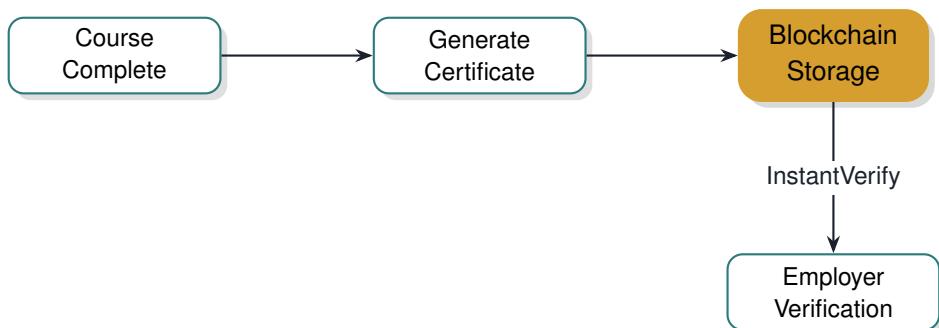


Figure 8.1: Blockchain-Based Credential System

8.4 Advanced AI Tutoring

Enhancing the AI assistant's capabilities:

- **Conversational Tutoring:** Enable back-and-forth dialogue where the AI guides students through problem-solving step-by-step
- **Voice Interaction:** Students can speak their questions naturally using speech-to-text

- **Multimodal Understanding:** Students can upload screenshots of problems they're stuck on, and the AI analyzes the image and provides help

8.5 Institutional Integration

Building connectors for existing educational tools:

- Single Sign-On (SSO) with university identity systems (SAML, OAuth)
- Grade book export to systems like Banner, PeopleSoft
- Calendar integration with Google Calendar, Outlook
- Library integration for embedding licensed educational resources

CHAPTER 9

Long-Term Vision (3-5 Years)



9.1 Personalized AI Agents

Moving beyond a shared AI assistant to personalized learning companions. Each student would have an AI agent that:

- Learns their specific learning style (visual vs. textual, fast vs. methodical pacing)
- Maintains long-term memory of what the student knows across all courses
- Proactively suggests content when detecting knowledge gaps
- Adapts explanations to the student's background and prior knowledge

9.2 VR/AR Immersive Learning

Leveraging virtual and augmented reality for experiential education:

- **Historical Immersion:** Walk through ancient Rome while learning about Roman Empire history
- **Medical Training:** Practice surgical procedures in a risk-free VR environment
- **Architecture Visualization:** Design buildings in CAD software and walk through them in VR
- **Language Learning:** Immerse yourself in a virtual Japanese marketplace, practicing conversations with AI-powered NPCs

9.3 Federated Learning for Privacy

Implementing machine learning that preserves student privacy. Federated learning allows models to train locally on each student's device, learning from behavior patterns without ever sending raw data to the cloud. Only model updates are shared and aggregated.

This enables highly personalized experiences while maintaining GDPR and FERPA compliance.

9.4 Global Education Network

Creating a marketplace where educators worldwide can collaborate:

- Cross-institutional course bundles (take Biology at University A, Chemistry at University B)
- Content licensing (University A licenses University B's excellent Machine Learning course)
- Shared certification standards recognized across institutions
- Translation services enabling courses to be taught in any language automatically

9.5 Neurofeedback Integration

Integrating biometric data to optimize learning. Using consumer EEG headbands or smart-watches, the system could detect:

- When students are confused or frustrated (adjust pacing or provide hints)
- When attention is waning (suggest a break)
- Optimal times of day for each student's learning based on cognitive performance patterns
- Stress levels during exams (provide calming exercises)

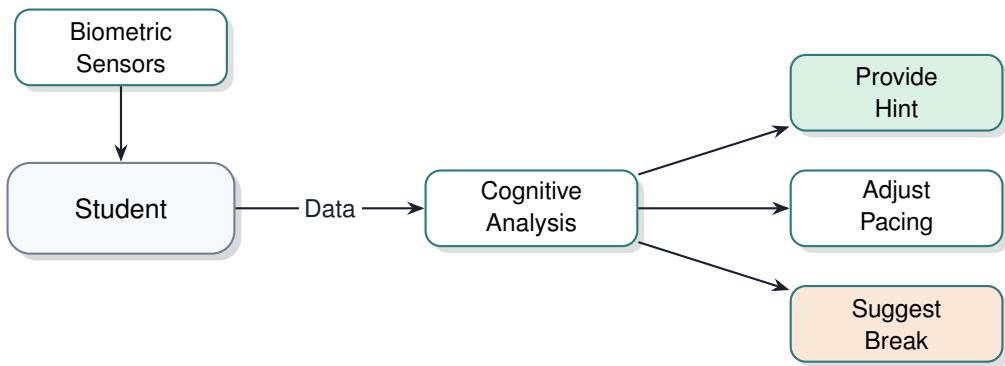
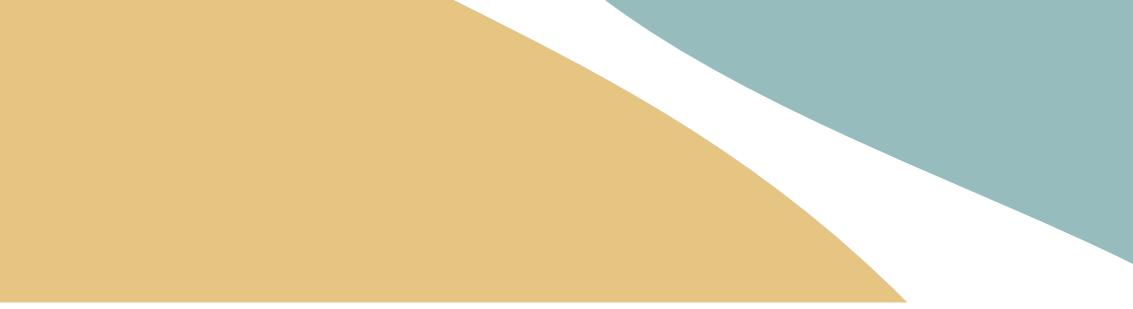


Figure 9.1: Neurofeedback-Driven Learning Optimization



CHAPTER 10

Conclusion



10.1 The Path Forward

Shikshak is more than a platform—it's a vision for what education can become when we thoughtfully apply cutting-edge technology to age-old challenges. Every architectural decision, every line of code, and every feature is guided by a simple question: *Does this help students learn better and teachers teach more effectively?*

The microservices architecture ensures Shikshak can scale from a small university pilot program to serving millions of learners globally without a complete rewrite. The AI-first approach ensures students always have support, regardless of time zones or instructor availability. The proctoring system restores trust in online credentials while respecting privacy.



10.2 Open Challenges

We acknowledge that perfect systems don't exist. Shikshak faces ongoing challenges:

- **AI Bias:** Ensuring language models don't perpetuate societal biases in their explanations
- **Access Inequality:** Students without reliable internet or modern devices may struggle with features like AI proctoring
- **Teacher Adaptation:** Many educators are skeptical of AI in education; we must prove value without threatening their role
- **Continuous Improvement:** Education evolves, technology evolves—Shikshak must evolve with them

10.3 Impact Potential

If successful, Shikshak could:

- Enable universities to offer high-quality online degrees at scale, expanding access to higher education
- Reduce teacher burnout by automating routine tasks, letting them focus on authentic human connection and mentorship
- Democratize elite education by making expert-created courses available globally at low cost
- Provide lifelong learning pathways as careers increasingly require continuous skill updates
- Generate datasets that advance learning science, revealing how humans acquire knowledge most effectively

Education is the great equalizer.

Technology like Shikshak has the potential to make high-quality learning accessible to anyone with an internet connection, regardless of geography, economic status, or institutional affiliation.

The future of education is intelligent, scalable, and deeply human.

Shikshak is building that future, one microservice at a time.

