

# Untitled 37

## RAG Service API Guide

This guide explains how to use the RAG service API endpoints for document ingestion and querying.

### Base URL

```
http://localhost:8000
```

For production, replace with your deployed service URL.

### Ingestion Endpoint

#### **POST /ingest**

Ingest documents (PDF, DOCX, TXT, or video files) into the RAG system.

#### Request Format

This endpoint accepts **multipart/form-data** with the following fields:

Field	Type	Required	Description
-------	------	----------	-------------

-----	-----	-----	-----
-------	-------	-------	-------

file   File   <input checked="" type="checkbox"/> Yes   The file to ingest (PDF, DOCX, TXT, MP4)
--

course_id   String   <input checked="" type="checkbox"/> Yes   Course identifier for multi-tenancy
--

module_id   String   <input checked="" type="checkbox"/> Yes   Module identifier within the course
--

source_type   String   <input checked="" type="checkbox"/> Yes   Source type: pdf , docx , txt , video , or notes
---

video_id   String   <input checked="" type="checkbox"/> No   Video ID (required if source_type is video )
---

notes_id   String   <input checked="" type="checkbox"/> No   Notes ID (required if source_type is notes )
---

#### Example: cURL

```
# Ingest a PDF file
```

```
curl -X POST "http://localhost:8000/ingest" \
-F "course_id=DL101" \
-F "module_id=M03" \
-F "source_type=pdf" \
-F "file=@lecture_notes.pdf"

# Ingest a video file

curl -X POST "http://localhost:8000/ingest" \
-F "course_id=DL101" \
-F "module_id=M02" \
-F "source_type=video" \
-F "video_id=video_123" \
-F "file=@lecture_video.mp4"
```

## Example: JavaScript (Fetch API)

```
const formData = new FormData();

formData.append('file', fileInput.files[0]);

formData.append('course_id', 'DL101');

formData.append('module_id', 'M03');

formData.append('source_type', 'pdf');

const response = await fetch('http://localhost:8000/ingest', {
  method: 'POST',
  body: formData
})
```

```
});  
  
const result = await response.json();  
  
console.log(result);
```

## Example: Python (Requests)

```
import requests  
  
  
  
files = {'file': open('lecture_notes.pdf', 'rb')}  
  
data = {  
  
    'course_id': 'DL101',  
  
    'module_id': 'M03',  
  
    'source_type': 'pdf'  
  
}  
  
  
  
response = requests.post('http://localhost:8000/ingest', files=files,  
data=data)  
  
print(response.json())
```

## Response

```
{  
  
    "job_id": "550e8400-e29b-41d4-a716-446655440000",  
  
    "status": "completed",  
  
    "message": "Successfully ingested 15 chunks",  
  
    "chunks_count": 15
```

```
}
```

## Response Fields:

- `job_id` : Unique identifier for the ingestion job
- `status` : Job status ( `queued` , `processing` , `completed` , or `failed` )
- `message` : Human-readable status message
- `chunks_count` : Number of text chunks created (only present for completed jobs)

## Query Endpoint

### **POST** /query

Query the RAG system with optional filtering by course and module.

#### Request Format

This endpoint accepts **JSON** payload:

Field	Type	Required	Default	Description
<code>query</code>	String	<input checked="" type="checkbox"/> Yes	-	The question to ask
<code>course_id</code>	String	<input type="checkbox"/> No	null	Filter results by specific course
<code>module_id</code>	String	<input type="checkbox"/> No	null	Filter results by specific module
<code>top_k</code>	Integer	<input type="checkbox"/> No	5	Number of chunks to retrieve (1-100)
<code>full_context</code>	Boolean	<input type="checkbox"/> No	false	If true, retrieve ALL chunks for the module
<code>include_sources</code>	Boolean	<input type="checkbox"/> No	true	If false, don't return source chunks

#### Example: cURL

```
# Basic query

curl -X POST "http://localhost:8000/query" \
-H "Content-Type: application/json" \
-d '{
```

```

    "query": "Explain the attention mechanism in transformers"

}'
```

# Query with course and module filtering

```

curl -X POST "http://localhost:8000/query" \
-H "Content-Type: application/json" \
-d '{
```

```

    "query": "Explain the attention mechanism",
    "course_id": "DL101",
    "module_id": "M03",
    "top_k": 5
```

```

}'
```

## Example: JavaScript (Fetch API)

```

const queryData = {

  query: "Explain the attention mechanism in transformers",
  course_id: "DL101",
  module_id: "M03",
  top_k: 5

};
```

const response = await fetch('http://localhost:8000/query', {
 method: 'POST',
 headers: {

```
'Content-Type': 'application/json'

},
body: JSON.stringify(queryData)

});

const result = await response.json();

console.log(result);
```

## Example: Python (Requests)

```
import requests

payload = {

    "query": "Explain the attention mechanism in transformers",

    "course_id": "DL101",

    "module_id": "M03",

    "top_k": 5

}

response = requests.post('http://localhost:8000/query', json=payload)

print(response.json())
```

## Response

```
{
    "answer": "The attention mechanism allows the model to focus on different parts of the input sequence when generating each output token. It computes
```

```

attention scores...",
"sources": [
{
  "chunk_id": "550e8400-e29b-41d4-a716-446655440000",
  "score": 0.89,
  "source_uri": "course/DL101/module/M03/lecture_notes.pdf",
  "source_type": "pdf",
  "text_preview": "Attention mechanism is a key component of transformer architecture...",
  "start_time_seconds": null,
  "end_time_seconds": null
},
],
"debug": {
  "search_latency_ms": 45.2,
  "llm_latency_ms": 1234.5,
  "total_latency_ms": 1279.7,
  "chunks_retrieved": 5,
  "tokens_used": 450,
  "cache_hit": false
}
}

```

## Response Fields:

- **answer** : The AI-generated answer based on retrieved context
- **sources** : List of source chunks used to generate the answer

- `chunk_id` : Unique identifier for the chunk
- `score` : Relevance score (0-1, higher is better)
- `source_uri` : Path/identifier of the source document
- `source_type` : Type of source ( `pdf` , `video` , `notes` , etc.)
- `text_preview` : Preview of the chunk text (max 200 chars)
- `start_time_seconds` / `end_time_seconds` : Time markers for video sources
  - `debug` : Performance and debugging information
- `search_latency_ms` : Time to retrieve relevant chunks
- `llm_latency_ms` : Time for LLM to generate answer
- `total_latency_ms` : Total query processing time
- `chunks_retrieved` : Number of chunks retrieved
- `tokens_used` : Estimated tokens used by LLM
- `cache_hit` : Whether the result was served from cache

## Health Check Endpoint

### **GET** /health

Check the health status of the RAG service and its dependencies.

### Example: cURL

```
curl http://localhost:8000/health
```

### Response

```
{  
  "status": "healthy",  
  "version": "1.0.0",  
  "dependencies": {  
    "llm": "stable",  
    "vector_db": "stable",  
    "indexer": "stable",  
    "cache": "stable"  
  }  
}
```

```
"qdrant_connected": true,  
"redis_connected": true  
}
```

---

## 🧪 Interactive API Documentation

FastAPI provides interactive API documentation at:

- **Swagger UI:** <http://localhost:8000/docs>
- **ReDoc:** <http://localhost:8000/redoc>

These interfaces allow you to test the API endpoints directly from your browser.

---

## 🔧 Common Use Cases

### 1. Ingest Course Materials

```
# Ingest lecture PDF  
  
curl -X POST "http://localhost:8000/ingest" \  
-F "course_id=CS101" \  
-F "module_id=intro" \  
-F "source_type=pdf" \  
-F "file=@intro_lecture.pdf"
```

```
# Ingest lecture video  
  
curl -X POST "http://localhost:8000/ingest" \  
-F "course_id=CS101" \  
-F "module_id=intro" \  
-F "source_type=video" \  
-F "file=@lecture_video.mp4"
```

```
-F "video_id=vid_001" \
-F "file=@intro_video.mp4"
```

## 2. Query Entire Course

```
curl -X POST "http://localhost:8000/query" \
-H "Content-Type: application/json" \
-d '{
  "query": "What are the main topics covered?",
  "course_id": "CS101"
}'
```

## 3. Query Specific Module

```
curl -X POST "http://localhost:8000/query" \
-H "Content-Type: application/json" \
-d '{
  "query": "Explain the key concepts",
  "course_id": "CS101",
  "module_id": "intro",
  "top_k": 10
}'
```

## 4. Get Full Module Context

```
curl -X POST "http://localhost:8000/query" \
```

```
-H "Content-Type: application/json" \
-d '{

    "query": "Summarize this module",

    "course_id": "CS101",

    "module_id": "intro",

    "full_context": true,

    "include_sources": false

}'
```

## ⚠ Error Handling

The API returns standard HTTP status codes:

- 200 : Success
- 400 : Bad Request (invalid parameters)
- 422 : Validation Error (missing required fields)
- 500 : Internal Server Error

Error response example:

```
{
    "detail": "Invalid source_type. Must be one of: ['pdf', 'docx', 'txt',
'vedeo', 'notes']"
}
```

## ❖ Python Helper Functions

Here are reusable Python functions for common operations:

### Ingest Notes (PDF, TXT, DOCX)

```
import requests

import os

import json


def ingest_notes(
    file_path: str,
    course_id: str,
    module_id: str,
    notes_id: str,
    base_url: str = "http://localhost:8000"
) -> dict:
    """
    Ingest a notes file (PDF, TXT, or DOCX) into the RAG system.

    Automatically detects file type from extension.
    """

    if not os.path.exists(file_path):
        raise FileNotFoundError(f"File not found: {file_path}")

    # Auto-detect source type from file extension
    _, ext = os.path.splitext(file_path)
    ext = ext.lower()

    ext_to_source_type = {".pdf": "pdf", ".txt": "txt", ".docx": "docx"}

    source_type = ext_to_source_type.get(ext)

    if not source_type:
        raise ValueError(f"Unsupported file type: {ext}")
```

```
with open(file_path, "rb") as f:

    files = {"file": (os.path.basename(file_path), f)}

    data = {

        "course_id": course_id,

        "module_id": module_id,

        "source_type": source_type,

        "notes_id": notes_id,

    }

    print(f"Ingesting {os.path.basename(file_path)}...")

    print(f" Course ID: {course_id}")

    print(f" Module ID: {module_id}")

    print(f" Source Type: {source_type}")

    response = requests.post(f"{base_url}/ingest", files=files,
data=data)

    result = response.json()

    print(json.dumps(result, indent=2))

    return result

# Usage

ingest_notes(

    file_path="lecture_notes.pdf",

    course_id="15",

    module_id="9",

    notes_id="4"

)
```

## Ingest Video

```
import requests

import os

import json


def ingest_video(
    file_path: str,
    course_id: str,
    module_id: str,
    video_id: str,
    base_url: str = "http://localhost:8000"
) -> dict:
    """
    Ingest a video file into the RAG system.

    Extracts audio and transcribes automatically.

    """
    if not os.path.exists(file_path):
        raise FileNotFoundError(f"File not found: {file_path}")

    with open(file_path, "rb") as f:
        files = {"file": (os.path.basename(file_path), f)}

        data = {
            "course_id": course_id,
            "module_id": module_id,
            "source_type": "video",
        }
```

```

        "video_id": video_id,
    }

    print(f"Ingesting {os.path.basename(file_path)}...")

    print(f" Course ID: {course_id}")

    print(f" Module ID: {module_id}")

    print(f" Video ID: {video_id}")

    print(" (This may take a while for audio transcription...)")

    response = requests.post(f"{base_url}/ingest", files=files,
data=data)

    result = response.json()

    print(json.dumps(result, indent=2))

    return result
}

# Usage

ingest_video(
    file_path="lecture_video.mp4",
    course_id="15",
    module_id="7",
    video_id="6"
)

```

## Query with Filters

```

import requests

import json

```

```
def query_rag(
    query: str,
    course_id: str = None,
    module_id: str = None,
    top_k: int = 10,
    full_context: bool = False,
    include_sources: bool = True,
    base_url: str = "http://localhost:8000"
) -> dict:
    """
    Query the RAG system with optional filtering.
    """

    payload = {
        "query": query,
        "top_k": top_k,
        "full_context": full_context,
        "include_sources": include_sources
    }

    if course_id:
        payload["course_id"] = course_id

    if module_id:
        payload["module_id"] = module_id

    response = requests.post(f"{base_url}/query", json=payload)

    result = response.json()

    print(json.dumps(result, indent=2))
```

```
    return result

# Query specific course and module

query_rag(
    query="What are the percentage of Sinhalese and Indian Tamils in Sri
Lanka?",

    course_id="15",

    module_id="7",

    top_k=10
)

# Query without module filter (entire course)

query_rag(
    query="Generate some questions on Sri Lanka?",

    course_id="15",

    top_k=10
)

# Get full context (all chunks) without source details in response

query_rag(
    query="Summarize the entire module",

    course_id="15",

    module_id="7",

    full_context=True,

    include_sources=False
```

)

---

## Notes

- **Multi-tenancy:** Use `course_id` and `module_id` to organize content hierarchically
- **File Formats:** Supported formats include PDF, DOCX, TXT, and MP4 (video)
- **Video Processing:** For videos, transcripts are extracted and chunked automatically (may take several minutes)
- **Caching:** Query results are cached in Redis for improved performance
- **Filters:** Combine `course_id` and `module_id` filters for precise content retrieval
- **Full Context Mode:** Use `full_context=True` to retrieve all chunks from a module instead of top-k semantic search
- **Source Toggle:** Set `include_sources=False` when you only need the answer and not the source chunks (useful with `full_context`)